

Algorithms for Massive Data Lecture Notes

Darshan Thaker

January 27, 2019

Contents

1	Approximate Counting	3
1.1	Morris's Algorithm	3
1.2	Variance Improvement by Averaging	5
1.3	Median Trick	7
2	Distinct Element Counting	9
2.1	Flajolet-Martin Algorithm	9

Lecture 1

Approximate Counting

We consider counting up to $n \in \mathbb{N}$ in the streaming model. To motivate this, consider a router that receives many requests and wants to count the number of requests. Naively, to store a counter for n , this would require $O(\log(n))$ bits of space (the big-O arises from the chosen base of the logarithm, which can be absorbed into a constant). Can we do better? The answer is a resounding no, *unless* we allow for approximation and randomness. If we allow approximation, the surprising answer is that we can use only $O(\log \log(n))$ bits of storage.

1.1 Morris's Algorithm

Algorithm 1.1, known as Morris's algorithm, is such an approximate algorithm.

Algorithm 1.1 Approximate Counting Algorithm [Morris 78]

Given: a stream of data.

Initialize variable $X = 0$.

Upon each new element in the stream, increment X with probability $\frac{1}{2^X}$, else do nothing.

Return estimator $E = 2^X - 1$.

We now analyze the correctness and performance of the Morris algorithm. First, let us start with the correctness of the algorithm. Let X_n be the value of X after n increments, and let $E_n = 2^{X_n} - 1$.

Theorem 1.1.1. *The estimator outputted by the Morris algorithm is a unbiased estimator, i.e.*

$$\mathbb{E}_{X_1, \dots, X_n}[E_n] = n$$

Proof. The proof proceeds by induction over n . Note that the expectation is over the n values of X (since these are all random variables). First, note the base case of $n = 0$. Clearly, $E_0 = 2^{X_0} - 1 = 2^0 - 1 = 0$, thus the base case is satisfied. Now, assume that the hypothesis holds for E_{n-1} , and we will show that it holds also for E_n . First, we will compute $\mathbb{E}_{X_1, \dots, X_n}[2^{X_n}]$. We can split this expectation into two expectations:

$$\mathbb{E}_{X_1, \dots, X_n}[2^{X_n}] = \mathbb{E}_{X_1, \dots, X_{n-1}}[\mathbb{E}_{X_n}[2^{X_n}]]$$

This trick is an important trick. It amounts to fixing the first $n - 1$ draws of X_1, \dots, X_{n-1} and then finding the expectation with respect to the last draw X_n . We can always do this by the rules of probability.

$$\begin{aligned}
&= \mathbb{E}_{X_1, \dots, X_{n-1}} [2^{X_{n-1}+1} \cdot \Pr[X_{n-1} \text{ incremented}] + 2^{X_{n-1}} \cdot \Pr[X_{n-1} \text{ not incremented}]] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}} \left[2^{X_{n-1}+1} \cdot \frac{1}{2^{X_{n-1}}} + 2^{X_{n-1}} \cdot \left(1 - \frac{1}{2^{X_{n-1}}}\right) \right] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}} [2 + 2^{X_{n-1}} - 1] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}} [2^{X_{n-1}}] + 1 \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}} [E_{n-1} + 1] + 1 \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}} [n - 1 + 1] + 1 \quad (\text{Inductive hypothesis}) \\
&= n + 1
\end{aligned}$$

Thus, by linearity of expectation (which we used above as well), we have that $\mathbb{E}_{X_1, \dots, X_n}[E_n] = \mathbb{E}_{X_1, \dots, X_n}[2^{X_n}] - 1 = n$, completing the proof. \square

We are not merely interested in the expectation of the estimator, but also in the concentration of measure of the estimator, which we now prove.

Theorem 1.1.2. *The estimator outputted by the Morris algorithm has the following property.*

$$\text{Var}[E_n] \leq \frac{3n(n-1)}{2} + 1 = O(n^2)$$

Proof. First, we note that instead of proving results about $\text{Var}[E_n]$, we can prove results about $\mathbb{E}[2^{2X_n}]$. To see this, note the following (where all expectations are taken over the random X_1, \dots, X_n):

$$\begin{aligned}
\text{Var}[E_n] &= \text{Var}[2^{X_n} - 1] \\
&= \text{Var}[2^{X_n}] \\
&= \mathbb{E}[(2^{X_n})^2] - (\mathbb{E}[2^{X_n}])^2 \\
&\leq \mathbb{E}[2^{2X_n}]
\end{aligned}$$

We now proceed via induction on n to show that $\mathbb{E}[2^{2X_n}] \leq \frac{3n(n-1)}{2} + 1$. Consider the base case for $n = 0$. We know that $\mathbb{E}[2^{2X_0}] = 1$ based on the initialization of the algorithm, which is less than or equal to $\frac{3 \cdot 0 \cdot (0-1)}{2} + 1$. Now, assume that the hypothesis holds for $\mathbb{E}_{X_1, \dots, X_{n-1}}[2^{2X_{n-1}}]$. Following the proof almost exactly as Theorem 1.1.1, we can conclude that:

$$\begin{aligned}
\mathbb{E}_{X_1, \dots, X_n}[2^{2X_n}] &= \mathbb{E}_{X_1, \dots, X_{n-1}}[\mathbb{E}_{X_n}[2^{2X_n}]] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}}\left[2^{2(X_{n-1}+1)} \cdot \frac{1}{2^{X_{n-1}}} + 2^{2X_{n-1}} \cdot \left(1 - \frac{1}{2^{X_{n-1}}}\right)\right] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}}[2^{X_{n-1}+2} + 2^{2X_{n-1}} - 2^{X_{n-1}}] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}}[4 \cdot 2^{X_{n-1}} + 2^{2X_{n-1}} - 2^{X_{n-1}}] \\
&= \mathbb{E}_{X_1, \dots, X_{n-1}}[3 \cdot 2^{X_{n-1}} + 2^{2X_{n-1}}] \\
&= 3 \cdot \mathbb{E}_{X_1, \dots, X_{n-1}}[2^{X_{n-1}}] + \mathbb{E}_{X_1, \dots, X_{n-1}}[2^{2X_{n-1}}] \\
&= 3 \cdot \mathbb{E}_{X_1, \dots, X_{n-1}}[E_{n-1} + 1] + \mathbb{E}_{X_1, \dots, X_{n-1}}[2^{2X_{n-1}}] \\
&= 3n + \mathbb{E}_{X_1, \dots, X_{n-1}}[2^{2X_{n-1}}] \quad (\text{Theorem 1.1.1}) \\
&\leq 3n + \frac{3(n-1)(n-2)}{2} \quad (\text{Inductive hypothesis}) \\
&= \frac{3n(n-1)}{2} + 1 \\
&= O(n^2)
\end{aligned}$$

By the principle of mathematical induction, we can conclude the claim holds for all natural numbers n . \square

Combining the above two theorems above, we can now state a probabilistic result about the approximation error of the estimator outputted by Morris's algorithm. Before stating this result, recall Chebyshev's inequality, which will come in handy many times. Chebyshev's inequality states that for any random variable X and for all $\lambda > 0$, we have the following bound:

$$\Pr[|X - \mathbb{E}[X]| \geq \lambda] \leq \frac{\text{Var}[X]}{\lambda^2}$$

Theorem 1.1.3. *The estimator outputted by Morris's algorithm E_n is in the range $n \pm O(n)$ with probability at least 90%.*

Proof. We use Chebyshev's inequality on the random variable E_n with the value of $\lambda^2 = 10\text{Var}[E_n]$. By Theorem 1.1.1 and Theorem 1.1.2, we know that $\mathbb{E}[E_n] = n$ and $\text{Var}[E_n] = O(n^2)$. Thus, by the Chebyshev bound, we have that E_n falls in the range $\mathbb{E}[E_n] \pm \lambda$ with probability at least 90%, or that E_n falls in the range $n \pm O(n)$ with probability at least 90%. \square

Finally, we observe that the Morris algorithm uses $O(\log \log(n))$ bits of storage with high probability, since the algorithm simply stores X and returns $E = 2^X - 1$. Since with high probability, E is "close" to n by the above theorem, we conclude that X requires $\log \log(n)$ bits of storage with high probability (at least 90%).

1.2 Variance Improvement by Averaging

We can improve the variance of the estimator outputted by the Morris algorithm using a generic technique. This leads to the so-called Morris+ algorithm, described in Algorithm 1.2, which simply outputs the average of k independent runs of the Morris algorithm.

Algorithm 1.2 Approximate Counting Algorithm with Better Variance [Morris 78]

Given: a stream of data, $k \in \mathbb{N}$

Run k independent copies of the Morris algorithm (Algorithm 1.1) to get k estimators, E_1, \dots, E_k .

Return estimator $E = \frac{1}{k} \sum_{i=1}^k E_i$.

We now analyze the correctness of this algorithm (using the same notation as the previous section).

Theorem 1.2.1. *The estimator outputted by the Morris+ algorithm is a unbiased estimator, i.e.*

$$\mathbb{E}_{E_1, \dots, E_n}[E] = n$$

Proof. This follows easily from the linearity of expectation and the application of Theorem 1.1.1. □

$$\begin{aligned} \mathbb{E}_{E_1, \dots, E_n}[E] &= \frac{1}{k} \sum_{i=1}^k \mathbb{E}_{E_i}[E_i] \\ &= \frac{1}{k} \cdot kn \\ &= n \end{aligned}$$

□

Similarly, we can analyze the variance of the estimator.

Theorem 1.2.2. *The estimator outputted by the Morris+ algorithm has the following property (for $n \geq 10$).*

$$\text{Var}[E] \leq \frac{3}{2} \cdot \frac{n^2}{k} = O\left(\frac{n^2}{k}\right)$$

Proof. By properties of the variance and the application of Theorem 1.1.2, we know the following:

$$\begin{aligned} \text{Var}\left[\frac{1}{k} \sum_{i=1}^k E_i\right] &= \frac{1}{k^2} \sum_{i=1}^k \text{Var}[E_i] \\ &\leq \frac{1}{k^2} \cdot k \cdot \frac{3}{2} n^2 && (\text{If } n \geq 10) \\ &= \frac{3}{2} \cdot \frac{n^2}{k} \\ &= O\left(\frac{n^2}{k}\right) \end{aligned}$$

□

Notice that the above theorem introduces an inverse dependence of the variance on k (which is a parameter we can pick). Thus, we can pick k to reduce the variance of the estimator as much as we desire. In particular, suppose we wish to achieve a $1 + \epsilon$ approximation to the desired output with high probability, i.e. the outputted estimator should be in the range $[n, (1 + \epsilon)n]$ with high probability. We can now state a result about this approximation using the Morris+ algorithm.

Theorem 1.2.3. *Using $k = O\left(\frac{1}{\epsilon^2}\right)$, the estimator outputted by the Morris+ algorithm E is in the range $n \pm \epsilon \cdot O(n)$ with probability at least 90%.*

Proof. We use Chebyshev's inequality on the random variable E with the value of $\lambda^2 = 10\text{Var}[E]$. By Theorem 1.2.1 and Theorem 1.2.2, we know that $\mathbb{E}[E] = n$ and $\text{Var}[E] = O\left(\frac{n^2}{k}\right)$. Thus, by the Chebyshev bound, we have that E falls in the range $\mathbb{E}[E] \pm \lambda$ with probability at least 90%, or that E falls in the range $n \pm \frac{1}{\sqrt{k}} O(n)$ with probability at least 90%. If $k = O\left(\frac{1}{\epsilon^2}\right)$, we have that E falls in the range $n \pm \epsilon \cdot O(n)$ with probability at least 90%. □

1.3 Median Trick

Observe that in the analysis in the previous sections, we used the Chebyshev bound to obtain a fixed confidence, i.e. with probability at least 90%, our estimator was ϵ -approximate. However, in many settings, we are interested in obtaining a specific confidence level via a parameter δ . Thus, with probability at least $1 - \delta$, we want our estimator to be ϵ -approximate (this should ring a bell with PAC learning). How can we boost our Morris algorithm to arbitrary confidence, given the Morris algorithm that works with 90% confidence? One solution is to use a Chebyshev bound with λ depending upon δ . We picked $\lambda^2 = 10\text{Var}[X]$ to achieve a 90% confidence. More generally, we can pick $\lambda^2 = \frac{1}{\delta}$, which gives $k = O\left(\frac{1}{\epsilon^2 \delta}\right)$ for the number of times to run the Morris algorithm to achieve a confidence level of at least $1 - \delta$. If δ is small enough, this dependence upon $\frac{1}{\delta}$ can be quite prohibitive, so a natural question is: can we do better? It turns out we can!

The answer is to use the *median trick*, which we will describe in detail below. Given an algorithm A that has a certain confidence level, or a specific probability that the output of the algorithm is in a correct range, we can apply the median trick to create an algorithm A' such that the probability that the output of the algorithm is in a correct range is at least $1 - \delta$. Additionally, we can avoid the $\frac{1}{\delta}$ dependence as in the solution above.

Algorithm 1.3 shows the median trick pseudocode. The algorithm is simple: simply run algorithm A k times and output the median of the estimators. Picking k wisely yields a very simple way to boost confidence.

Algorithm 1.3 Median Trick

Given: Algorithm A with a certain confidence level, $k \in \mathbb{N}$

Run k independent copies of Algorithm A to get k estimators, E_1, \dots, E_k .

Return estimator E which is the median of E_1, \dots, E_k .

Before we proceed with the analysis, let us remind ourselves of the Chernoff and Hoeffding bounds. Suppose X_1, \dots, X_n are iid random variables that take values in $[0, 1]$. Let $\mu = \mathbb{E}[\sum_{i=1}^n X_i]$, and ϵ be a value in $[0, 1]$. Then, the Chernoff/Hoeffding bounds state:

$$\Pr \left[\left| \sum_{i=1}^n X_i - \mu \right| > \epsilon \mu \right] \leq 2 \exp \left\{ -\frac{\epsilon^2 \mu}{3} \right\}$$

These bounds give concentration inequalities on how far a sum of iid random variables will deviate from its expectation. We are now ready to prove a result about the median trick. Note that the median trick intuitively will only work for boosting confidence levels that are strictly greater than 0.5. This should make sense because a confidence level of 0.5 can be achieved by random guessing, so it should make sense that we cannot extract any useful information from random guessing.

Theorem 1.3.1. *If we take $k = O\left(\log \frac{1}{\delta}\right)$, the median trick will boost the confidence of Algorithm A from a confidence level c to at least $1 - \delta$.*

Proof. Let A_1, \dots, A_k be the output of Algorithm A on the k instances it is run. Let X_i be an indicator variable that is 1 when A_i is "correct" (in the desired approximate range), and 0 otherwise. Let $\mu = \mathbb{E}[\sum_{i=1}^k X_i]$. Then, $\mathbb{E}[X_i] \geq \Pr[X_i = 1] = c$ (because c is the confidence level of Algorithm A). Thus, we conclude that $\mu \geq c \cdot k$ (by independence). Note that Algorithm A' is correct if greater than half of the k estimators are correct, since the outputted value is the median. Note that this is a sufficient, not necessary condition. Because each X_i is an indicator variable, this condition is equivalent to the condition that $\sum_{i=1}^k X_i \geq 0.5k$. We would like the probability of this event to be at least $1 - \delta$. Instead, we solve the easier problem of solving for k such that the probability (of failure) that $\sum_{i=1}^k X_i < 0.5k$ is upper bounded by δ . Observe the algebraic manipulation of this condition below to allow us to use the Chernoff bound:

$$\begin{aligned}
\Pr \left[\sum_{i=1}^k X_i < 0.5k \right] &= \Pr \left[\sum_{i=1}^k X_i - ck < (0.5 - c) \cdot k \right] \\
&\leq \Pr \left[\left| \sum_{i=1}^k X_i - ck \right| > (c - 0.5) \cdot k \right] && (c > 0.5) \\
&\leq \Pr \left[\left| \sum_{i=1}^k X_i - \mu \right| > \frac{c - 0.5}{c} \cdot \mu \right] \\
&\leq 2 \exp \left\{ -\frac{1}{3} \left(\frac{c - 0.5}{c} \right)^2 \mu \right\} && (\text{Chernoff bound}) \\
&\leq 2 \exp \left\{ -\frac{1}{3} \left(\frac{c - 0.5}{c} \right)^2 ck \right\} && (\mu \geq ck)
\end{aligned}$$

We set this probability to be less than δ and solve for a suitable value of k :

$$\begin{aligned}
2 \exp \left\{ -\frac{1}{3} \left(\frac{c - 0.5}{c} \right)^2 ck \right\} &< \delta \\
\frac{1}{3} \frac{(c - 0.5)^2}{c} k &> \log \left(\frac{2}{\delta} \right) \\
k &> 3 \frac{c}{(c - 0.5)^2} \log \left(\frac{2}{\delta} \right) \\
k &> O \left(\log \left(\frac{1}{\delta} \right) \right)
\end{aligned}$$

□

With the median trick, we can improve Morris's algorithm to use $k = O \left(\frac{1}{\epsilon^2} \log \frac{1}{\delta} \right)$ to achieve an ϵ -approximate output with confidence $1 - \delta$. The algorithm runs the Morris+ algorithm first to achieve a fixed confidence level and then uses the median trick to boost the confidence level.

Lecture 2

Distinct Element Counting

We turn to the problem of counting the number of distinct elements in a stream. Specifically, consider a stream of m numbers x_1, \dots, x_m , where each $x_i \in [n]$. We want to construct an algorithm to output the number of distinct elements seen in the stream. There are several naive algorithms that should immediately come to mind. First, we can simply store a bit array of length n (requiring $O(n)$ bits of storage) and keep track of elements seen so far. Similarly, we can store all the input elements (requiring $O(m \log n)$ bits of storage). However, if m or n is very large, both these algorithms can be infeasible. We now turn to the Flajolet-Martin algorithm for distinct element counting, which uses roughly $O(\log n)$ bits of storage.

2.1 Flajolet-Martin Algorithm

The idea behind the Flajolet-Martin algorithm is to make use of a specific hash function oracle. Suppose that we were given access to a random hash function $h : [n] \rightarrow [0, 1]$. Specifically, for any number in 1 to n , the hash function would give a real number from 0 to 1, and such a hash function can be picked once at the beginning of the algorithm and queries to this function can be made afterwards. For now, we simply suppose that such a hash function exists and we can efficiently store it (and we will prove this later). Given this hash function h , the Flajolet-Martin (FM) algorithm is summarized in Algorithm 2.4.

Algorithm 2.4 Flajolet-Martin (FM) Algorithm for Distinct Element Counting

Given: Random hash function $h : [n] \rightarrow [0, 1]$

Initialize variable $Z = 1$.

Upon seeing an element $X = i$, set $Z = \min(Z, h(i))$.

Return estimator $E = \frac{1}{Z} - 1$.

The intuition behind this algorithm is the counter we store, Z , will store the smallest hashed value seen in the stream. Notice that for duplicate elements in the stream, the value of Z does not change because of the static hash function h . The minimum hashed value seen should, in expectation, store some information about the number of elements seen. This is because, in expectation, a random hash function will uniformly distribute each element in the domain over $[0, 1]$. This will be formalized in more detail below. We now analyze the FM algorithm. Let d be the true number of distinct elements in the stream.

Theorem 2.1.1. *The counter Z in by the Flajolet-Martin algorithm has the following property:*

$$\mathbb{E}[Z] = \frac{1}{d+1}$$

Proof. Observe that duplicate elements in the stream do not affect the value of Z . Thus, Z is the minimum of d random variables lying in the range $[0, 1]$. Pick any $a \in [0, 1]$ at random. The probability that a is less than Z is $\frac{1}{d+1}$, since this is the probability that a is the smallest of $d+1$ numbers distributed independently over the interval $[0, 1]$ (the probability that the $d+1$ numbers are all equal is 0 by properties of a continuous density). Note that this probability is also exactly equal to $\mathbb{E}[Z]$, since $\mathbb{E}[Z]$ is the expected value of the minimum of d random variables in $[0, 1]$. Equating these two quantities, we have that $\mathbb{E}[Z] = \frac{1}{d+1}$. \square

Theorem 2.1.2.

$$\text{Var}[Z] \leq \frac{2}{d^2}$$

Proof.

\square